

## 1 Assorted Joins

- Companies: (company\_id, industry, ipo\_date)
- NYSE: (company\_id, date, trade, quantity)

We have 20 pages of memory, and we want to join two tables Companies and NYSE on C.company\_id = N.company\_id. Attribute company\_id is the primary key for Companies. For every tuple in Companies, assume there are 4 matching tuples in NYSE.

NYSE contains [N] = 100 pages, NYSE holds pN = 100 tuples per page.

Companies contains [C] = 50 pages, C holds pC = 50 tuples per page.

There are unclustered B+ tree indexes of height 1 on C.company\_id and N.company\_id.

(a) How many disk I/Os are needed to perform a simple nested loops join?

$$[C] + pC * [C] * [N] = 50 + 50 * 50 * 100 = 250050$$

(b) How many disk I/Os are needed to perform a block nested loops join?

With N as the outer relation:  $[N] + \text{ceil}([N]/B-2) * [C] = 100 + \text{ceil}(100/18) * 50 = 400$  I/Os

With C as the outer relation:  $[C] + \text{ceil}([C]/B-2) * [N] = 50 + \text{ceil}(50/18) * 100 = 350$  I/Os

Most of the time, putting the smaller relation on the outside leads to a better I/O cost for BNLJ, but this is not always true. Therefore, we try both options, computing the I/O cost with N on the outside and with C on the outside and taking the smaller of the two.

I/O cost for BNLJ: 350 I/Os

(c) How many disk I/Os are needed to perform an index nested loops join?

$$[C] + [C] * pC * (\text{cost to find matching NYSE tuples}) = 50 + 50 * 50 * (2 + 4) = 15,050$$

You should check which relation on the outside results in a lower I/O cost. See 1(b) for more description.

(d) For this part only, assume the index on NYSE.company\_id is clustered. What is the cost of an index nested loops join using companies as the outer relation?

$$[C] + [C] * pC * (\text{cost to find matching NYSE tuples}) = 50 + 50 * 50 * (2 + \text{ceil}(4/100)) = 7,550 \text{ I/Os}$$

You should check which relation on the outside results in a lower I/O cost. See 1(b) for more description.

- (e) How many disk I/Os are needed to perform a sort merge join without optimization? If we can perform the sort merge join optimization, how many disk I/Os are needed with optimization?

**Without Optimization:**

Sorting N:

Pass 0 -  $\text{ceil}(100/20) = 5$  sorted runs of 20 pages each

Pass 1 -  $\text{ceil}(5/19) = 1$  sorted run of 100 pages each

Total I/Os:  $4 * (100 \text{ pages}) = 400 \text{ I/Os}$

Sorting C:

Pass 0 -  $\text{ceil}(50/20) = 3$  sorted runs of 20 pages, 20 pages, and 10 pages

Pass 1 -  $\text{ceil}(3/19) = 1$  sorted run of 50 pages

Total I/Os:  $4 * (50 \text{ pages}) = 200 \text{ I/Os}$

Joining:  $[C] + [N] = 150 \text{ I/Os}$

Total:  $200 + 400 + 150 = 750 \text{ I/Os}$

**With Optimization:**

During the 2nd to last pass, we produce 5 sorted runs of N and 3 sorted runs of C. Since the number of runs of C + the number of runs of N  $\leq 20 - 1$ , we can optimize sort merge join and combine the last sorting pass and final merging pass to save  $2 * ([C] + [N]) \text{ I/Os}$ .

Total I/Os =  $750 - 2(50+100) = 450 \text{ I/Os}$

- (f) How many disk I/Os are needed to perform a hash join? Assume uniform partitioning.

No recursive partitioning required.

Partitioning phase:

$\text{ceil}([N]/(B - 1)) = 6$  pages per partition for N, 19(6) pages

$\text{ceil}([C]/(B - 1)) = 3$  pages per partition for C, 19(3) pages

Partitioning Phase: 100 I/Os to read for N + 19(6) I/Os to write for N + 50 I/Os to read for C + 19(3) I/Os to write for C = 321 I/Os

Build and Probe:  $19(6) + 19(3) = 171 \text{ I/Os}$  to read for N and C

Total:  $321 + 171 = \mathbf{492 \text{ I/Os}}$

## 2 Grace Hash Join

We have 2 tables – Catalog and Transactions.

Catalog has a total of 100 pages and 20 tuples per page. Transactions has a total of 50 pages and 50 tuples per page. Assume the hash functions uniformly distribute the data for both tables.

- (a) If we had 10 buffer pages, how many partitioning phases would we require for grace hash join? Consider which table we should build the hash table in the probing phase on.

T is smaller, so we need its partitions to be at most  $B - 2 = 8$  pages. After 1 partitioning pass, we have partitions of size 6, which is  $\leq 8$  so we only need **1 partitioning pass**.

(b) What is the I/O cost for the grace hash join then?

We need 1 partitioning pass.

Partitioning phase:

$\text{ceil}([C]/(B - 1)) = 12$  pages per partition for C, 12(9) pages in total after partitioning

$\text{ceil}([T]/(B - 1)) = 6$  pages per partition for T, 6(9) pages in total after partitioning

Partitioning IOs: 100 I/Os to read from Catalog + 12(9) to write for Catalog + 50 I/Os to read from Transactions + 6(9) to write for Transactions = 312 I/Os

Probing phase: 12(9) + 6(9) = 162 I/Os to read from Catalog and Transactions

Total: 312 + 162 = **474 I/Os**

(c) For the above question, if we only had 8 buffer pages, how many partitioning phases would there be?

T is smaller, so we need its partitions to be at most  $B - 2 = 6$  pages. After 1 partitioning pass, we have partitions of size 8, which is too big to fit in B-2 buffer pages. We need a second partitioning pass.  $8 / 7 = 1.1 \rightarrow 2$  pages, which is small enough to fit in B-2 buffer pages. Therefore, we need **2 passes** in total.

(d) What will be the I/O cost?

Partitioning phase:

$\text{ceil}([C]/(B - 1)) = 15$  pages per partition for C

$\text{ceil}([T]/(B - 1)) = 8$  pages per partition for T

$\text{ceil}([C]/(B - 1)) = 3$  pages per partition for second pass for C

$\text{ceil}([T]/(B - 1)) = 2$  pages per partition for second pass for T

Partitioning IOs: [100 + 50] (1st read) + [15(7) + 8(7)] (1st write) + [15(7) + 8(7)] (2nd read) + [3(49) + 2(49)] (2nd write) = 717 I/Os

Build and Probe Phase: 3(49) + 2(49) = 245 IOs

Total: 717 + 245 = **962 I/Os**