

# SQL II

R & G - Chapter 5



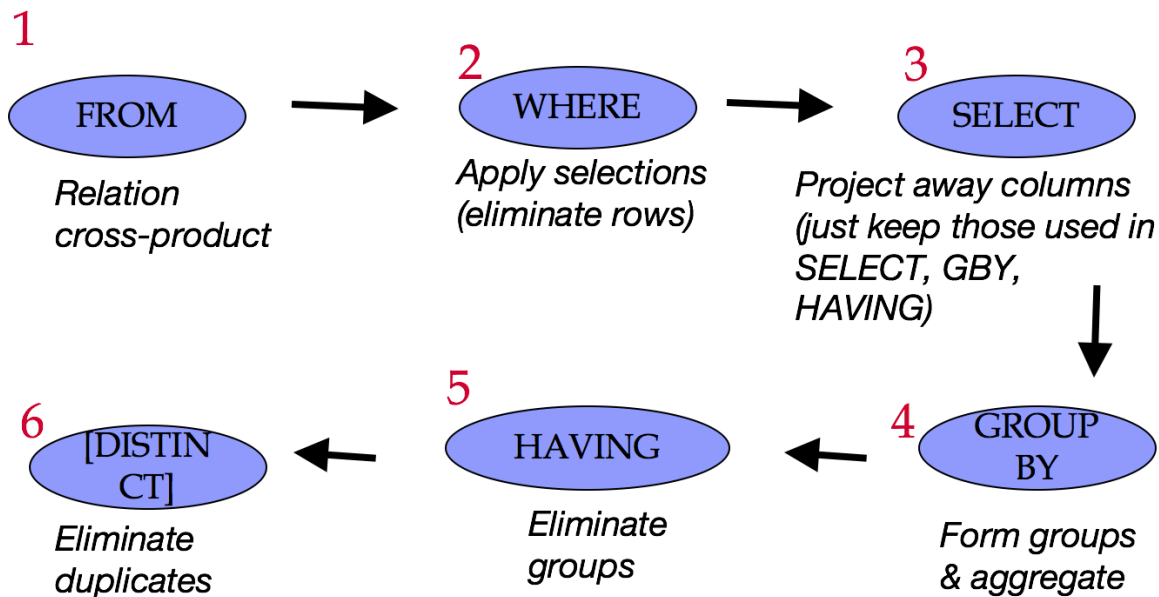
# SQL DML 1:

## Basic Single-Table Queries



- **SELECT** [**DISTINCT**] *<column expression list>*  
**FROM** *<single table>*  
[**WHERE** *<predicate>*]  
[**GROUP BY** *<column list>*  
[**HAVING** *<predicate>*] ]  
[**ORDER BY** *<column list>*]  
[**LIMIT** *<integer>*];

# Conceptual SQL Evaluation



SELECT      [DISTINCT] *target-list*  
FROM        *relation-list*  
WHERE       *qualification*  
GROUP BY   *grouping-list*  
HAVING     *group-qualificati*

# Putting it all together



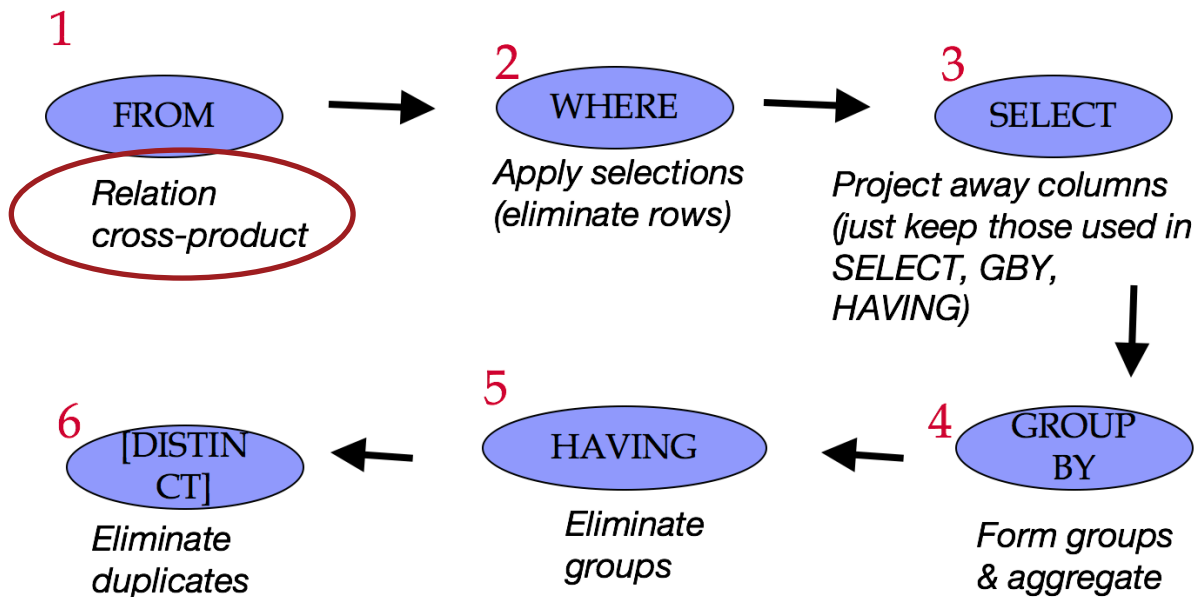
- **SELECT** S.dept, **AVG**(S.gpa), **COUNT**(\*)  
**FROM** Students S  
**WHERE** S.gender = 'F'  
**GROUP BY** S.dept  
**HAVING COUNT**(\*) >= 2  
**ORDER BY** S.dept;

# Join Queries



- SELECT [DISTINCT] *<column expression list>*  
**FROM** *<table1 [AS t1], ... , tableN [AS tn]>*  
[WHERE *<predicate>*]  
[GROUP BY *<column list>*[HAVING *<predicate>*] ]  
[ORDER BY *<column list>*];

# Conceptual SQL Evaluation, cont



SELECT [DISTINCT] target-list  
FROM relation-list  
WHERE qualification  
GROUP BY grouping-list  
HAVING group-qualificati

# Cross (Cartesian) Product



- All pairs of tuples, concatenated

## Sailors

sid	sname	rating	age
1	Popeye	10	22
2	OliveOyl	11	39
3	Garfield	1	27
4	Bob	5	19

## Reserves

sid	bid	day
1	102	9/12
2	102	9/13
1	101	10/01

sid	sname	rating	age	sid	bid	day
1	Popeye	10	22	1	102	9/12
1	Popeye	10	22	2	102	9/13
1	Popeye	10	22	1	101	10/01
2	OliveOyl	11	39	1	102	9/12
...	...	...	...	...	...	...

Slide Deck Title

# Find sailors who've reserved a boat



```
SELECT S.sid
FROM Sailors AS S, Reserves AS R
WHERE S.sid=R.sid
```

sid	sname	rating	age
1	Popeye	10	22
2	OliveOyl	11	39
3	Garfield	1	27
4	Bob	5	19

sid	bid	day
1	102	9/12
2	102	9/13
1	101	10/01

sid	sname	rating	age	sid	bid	day
1	Popeye	10	22	1	102	9/12
1	Popeye	10	22	2	102	9/13
1	Popeye	10	22	1	101	10/01
2	OliveOyl	11	39	1	102	9/12
...	...	...	...	...	...	...

Slide Deck Title



# Find sailors who've reserved a boat cont



```
SELECT S.sid
FROM Sailors AS S, Reserves AS R
WHERE S.sid=R.sid
```

sid	sname	rating	age
1	Popeye	10	22
2	OliveOyl	11	39
3	Garfield	1	27
4	Bob	5	19

sid	bid	day
1	102	9/12
2	102	9/13
1	101	10/01

sid	sname	bid
1	Popeye	102
1	Popeye	101
2	OliveOyl	102

# Column Names and Table Aliases



```
SELECT Sailors.sid, sname, bid  
FROM Sailors, Reserves  
WHERE Sailors.sid = Reserves.sid
```

```
SELECT S.sid, sname, bid  
FROM Sailors AS S, Reserves AS R  
WHERE S.sid = R.sid
```

# More Aliases



```
SELECT x.sname, x.age,  
       y.sname AS sname2,  
       y.age AS age2  
FROM Sailors AS x, Sailors AS y  
WHERE x.age > y.age
```

sname	age	sname2	age2
Popeye	22	Bob	19
OliveOyl	39	Popeye	22
OliveOyl	39	Garfield	27
OliveOyl	39	Bob	19
Garfield	27	Popeye	22
Garfield	27	Bob	19

- Table aliases in the FROM clause
  - Needed when the same table used multiple times (“self-join”)
- Column aliases in the SELECT clause

# Arithmetic Expressions



- **SELECT S.age, S.age-5 AS age1, 2\*S.age AS age2**  
FROM Sailors AS S  
WHERE S.sname = 'Popeye'
  
- **SELECT S1.sname AS name1, S2.sname AS name2**  
FROM Sailors AS S1, Sailors AS S2  
WHERE **2\*S1.rating = S2.rating - 1**

# SQL Calculator!



SELECT

log(1000) as three,

exp(ln(2)) as two,

cos(0) as one,

ln(2\*3) = ln(2) + ln(3) as sanity;

# String Comparisons



- Old School SQL  
SELECT S.sname  
FROM Sailors S  
WHERE **S.sname LIKE 'B\_%'**
- Standard Regular Expressions  
SELECT S.sname  
FROM Sailors S  
WHERE **S.sname ~ 'B.\*'**

# Combining Predicates



- Subtle connections between:
  - Boolean logic in WHERE (i.e., AND, OR)
  - Traditional Set operations (i.e. INTERSECT, UNION)
- Let's see some examples...

Sid's of sailors who reserved a red **OR** a green boat



```
SELECT R.sid  
FROM   Boats B, Reserves R  
WHERE  R.bid=B.bid AND  
        (B.color='red' OR B.color='green')
```



# Sid's of sailors who reserved a red **OR** a green boat Pt 2



```
SELECT R.sid
FROM Boats B,Reserves R
WHERE R.bid=B.bid AND
      (B.color='red' OR B.color='green')
```

**VS...**

```
SELECT R.sid
FROM Boats B, Reserves R
WHERE R.bid=B.bid AND B.color='red'
```

UNION ALL

```
SELECT R.sid
FROM Boats B, Reserves R
WHERE R.bid=B.bid AND B.color='green'
```

# Sid's of sailors who reserved a red **OR** a green boat Pt 3



```
SELECT R.sid
FROM Boats B,Reserves R
WHERE R.bid=B.bid AND
      (B.color='red' AND B.color='green')
```

**VS...**

```
SELECT R.sid
FROM Boats B, Reserves R
WHERE R.bid=B.bid AND B.color='red'
```

INTERSECT

```
SELECT R.sid
FROM Boats B, Reserves R
WHERE R.bid=B.bid AND B.color='green'
```

# Find sailors who have **not** reserved a boat



```
SELECT S.sid  
FROM Sailors S
```

```
EXCEPT
```

```
SELECT S.sid  
FROM Sailors S, Reserves R  
WHERE S.sid=R.sid
```

# Set Semantics



- Set: a collection of distinct elements
- Standard ways of manipulating/combining sets
  - Union
  - Intersect
  - Except
- Treat tuples within a relation as elements of a set

# Default: Set Semantics

$R = \{A, A, A, A, B, B, C, D\}$

$S = \{A, A, B, B, B, C, E\}$

- UNION

$\{A, B, C, D, E\}$

- INTERSECT

$\{A, B, C\}$

- EXCEPT

$\{D\}$

Note: Think of each letter as being a **tuple** in **relation**.

ex:

**A:** (Jim, 18, English, 4.0)

**B:** (Marcela, 20, CS, 3.8)

**C:** (Gail, 19, Statistics, 3.74)

**D:** (Goddard, 20, Math, 3.8)

# “ALL”: Multiset Semantics

$R = \{A, A, A, A, B, B, C, D\} = \{A(4), B(2), C(1), D(1)\}$

$S = \{A, A, B, B, B, C, E\} = \{A(2), B(3), C(1), E(1)\}$

# “UNION ALL”: Multiset Semantics

$R = \{A, A, A, A, B, B, C, D\} = \{A(4), B(2), C(1), D(1)\}$

$S = \{A, A, B, B, B, C, E\} = \{A(2), B(3), C(1), E(1)\}$

- UNION ALL: sum of cardinalities

$\{A(4+2), B(2+3), C(1+1), D(1+0), E(0+1)\}$

$= \{A, A, A, A, A, A, B, B, B, B, B, C, C, D, E\}$

# “INTERSECT ALL”: Multiset Semantics

$R = \{A, A, A, A, B, B, C, D\} = \{A(4), B(2), C(1), D(1)\}$

$S = \{A, A, B, B, B, C, E\} = \{A(2), B(3), C(1), E(1)\}$

- INTERSECT ALL: min of cardinalities  
 $\{A(\min(4,2)), B(\min(2,3)), C(\min(1,1)),$   
 $D(\min(1,0)), E(\min(0,1))\}$   
 $= \{A, A, B, B, C\}$



# “EXCEPT ALL”: Multiset Semantics

$R = \{A, A, A, A, B, B, C, D\} = \{A(4), B(2), C(1), D(1)\}$

$S = \{A, A, B, B, B, C, E\} = \{A(2), B(3), C(1), E(1)\}$


- EXCEPT ALL: difference of cardinalities  
 $\{A(4-2), B(2-3), C(1-1), D(1-0), E(0-1)\}$   
 $= \{A, A, D, \}$

# Nested Queries: IN

- *Names of sailors who've reserved boat #102:*

```
SELECT S.sname  
FROM   Sailors S  
WHERE  S.sid IN  
      (SELECT R.sid  
       FROM   Reserves R  
       WHERE  R.bid=102)
```

**subquery**



# Nested Queries: NOT IN

- *Names of sailors who've not reserved boat #103:*

```
SELECT S.sname  
FROM Sailors S  
WHERE S.sid NOT IN  
      (SELECT R.sid  
       FROM Reserves R  
       WHERE R.bid=103)
```

# Nested Queries: EXISTS

- *This is a bit odd, but it is legal:*

```
SELECT S.sname  
FROM Sailors S  
WHERE EXISTS  
    (SELECT R.sid  
     FROM Reserves R  
     WHERE R.bid=103)
```

# Nested Queries with Correlation

- *Names of sailors who've reserved boat #102:*

```
SELECT S.sname
FROM   Sailors S
WHERE EXISTS
      (SELECT *
       FROM Reserves R
       WHERE R.bid=102 AND S.sid=R.sid)
```

- Correlated subquery is recomputed for each Sailors tuple.

# More on Set-Comparison Operators

- We've seen: IN, EXISTS
- Can also have: NOT IN, NOT EXISTS
- Other forms: op ANY, op ALL

Find sailors whose rating is greater than that of *some* sailor called Popeye:

```
SELECT *  
FROM   Sailors S  
WHERE  S.rating > ANY  
        (SELECT S2.rating  
         FROM   Sailors S2  
         WHERE  S2.sname='Popeye')
```

# A Tough One: “Division”

- Relational Division: “Find sailors who’ve reserved all boats.”  
Said differently: “sailors with no counterexample missing boats”

```
SELECT S.sname
FROM Sailors S
WHERE NOT EXISTS
  (SELECT B.bid
   FROM Boats B
   WHERE NOT EXISTS (SELECT R.bid
                     FROM Reserves R
                     WHERE R.bid=B.bid
                     AND R.sid=S.sid ))
```

# ARGMAX? Pt 1

- The sailor with the highest rating
- Correct or Incorrect?

```
SELECT MAX(S.rating)  
FROM Sailors S;
```

**VS**

```
SELECT S.*, MAX(S.rating)  
FROM Sailors S;
```



# ARGMAX? Pt 2

- The sailor with the highest rating
- Correct or Incorrect? Same or different?

```
SELECT *  
FROM Sailors S  
WHERE S.rating >= ALL  
      (SELECT S2.rating  
       FROM Sailors S2)
```

**VS**

```
SELECT *  
FROM Sailors S  
WHERE S.rating =  
      (SELECT MAX(S2.rating)  
       FROM Sailors S2)
```

# ARGMAX? Pt 3

- The sailor with the highest rating
- Correct or Incorrect? Same or different?

```
SELECT *  
FROM Sailors S  
WHERE S.rating >= ALL  
      (SELECT S2.rating  
       FROM Sailors S2)
```

**VS**

```
SELECT *  
FROM Sailors S  
ORDER BY rating DESC  
LIMIT 1;
```

# “Inner” Joins: Another Syntax

```
SELECT s.*, r.bid  
FROM Sailors s, Reserves r  
WHERE s.sid = r.sid  
AND ...
```

```
SELECT s.*, r.bid  
FROM Sailors s INNER JOIN Reserves r  
ON s.sid = r.sid  
WHERE ...
```

# Join Variants

```
SELECT <column expression list>
FROM table_name
    [INNER | NATURAL
     | {LEFT | RIGHT | FULL } {OUTER}] JOIN table_name
    ON <qualification_list>
WHERE ...
```

- **INNER** is default
- Inner join what we've learned so far
  - Same thing, just with different syntax.

# Inner/Natural Joins

```
SELECT s.sid, s.sname, r.bid  
FROM Sailors s, Reserves r  
WHERE s.sid = r.sid  
      AND s.age > 20;
```

```
SELECT s.sid, s.sname, r.bid  
FROM Sailors s INNER JOIN Reserves r  
ON s.sid = r.sid  
WHERE s.age > 20;
```

```
SELECT s.sid, s.sname, r.bid  
FROM Sailors s NATURAL JOIN Reserves r  
WHERE s.age > 20;
```

- **ALL 3 ARE EQUIVALENT!**
- “NATURAL” means equi-join for pairs of attributes with the same name

# Left Outer Join

- Returns all matched rows, and preserves all unmatched rows from the table on the left of the join clause
  - (use nulls in fields of non-matching tuples)

```
SELECT s.sid, s.sname, r.bid  
FROM Sailors2 s LEFT OUTER JOIN Reserves2 r  
ON s.sid = r.sid;
```

Returns all sailors & bid for boat in any  
of their reservations

Note: no match for s.sid? r.bid IS NULL!

# Right Outer Join

- Returns all matched rows, and preserves all unmatched rows from the table on the right of the join clause
  - (use nulls in fields of non-matching tuples)

```
SELECT r.sid, b.bid, b.bname  
FROM Reserves2 r RIGHT OUTER JOIN Boats2 b  
ON r.bid = b.bid
```

Returns all boats and sid for any sailor associated with the reservation.

Note: no match for b.bid? r.sid IS NULL!

# Full Outer Join

- Returns all (matched or unmatched) rows from the tables on both sides of the join clause

```
SELECT r.sid, b.bid, b.bname  
FROM Reserves2 r FULL OUTER JOIN Boats2 b  
ON r.bid = b.bid
```

- Returns all boats & all information on reservations
- No match for r.bid?
  - b.bid IS NULL AND b.bname IS NULL!
- No match for b.bid?
  - r.sid IS NULL!



# Views: Named Queries

**CREATE VIEW** *view\_name*  
*AS select\_statement*

- Makes development simpler
- Often used for security
- Not “materialized”

CREATE VIEW Redcount

AS SELECT B.bid, COUNT(\*) AS scount  
FROM Boats2 B, Reserves2 R  
WHERE R.bid=B.bid AND B.color='red'  
GROUP BY B.bid

# Views Instead of Relations in Queries

```
CREATE VIEW Redcount
AS SELECT B.bid, COUNT(*) AS scout
   FROM Boats2 B, Reserves2 R
  WHERE R.bid=B.bid AND B.color='red'
 GROUP BY B.bid;
```

```
SELECT * from redcount;
```

bid	scount
102	1

```
SELECT bname, scout
FROM Redcount R, Boats2 B
WHERE R.bid=B.bid
AND scout < 10;
```

# Subqueries in FROM

Like a “view on the fly”!

```
SELECT bname, scout  
FROM Boats2 B,  
(SELECT B.bid, COUNT (*)  
    FROM Boats2 B, Reserves2 R  
    WHERE R.bid = B.bid AND B.color = 'red'  
    GROUP BY B.bid) AS Reds(bid, scout)  
  
WHERE Reds.bid=B.bid  
    AND scout < 10
```

# WITH a.k.a. common table expression (CTE)

## Another “view on the fly” syntax:

```
WITH Reds(bid, scout) AS  
(SELECT B.bid, COUNT (*)  
FROM Boats2 B, Reserves2 R  
WHERE R.bid = B.bid AND B.color = 'red'  
GROUP BY B.bid)
```

```
SELECT bname, scout  
FROM Boats2 B, Reds  
WHERE Reds.bid=B.bid  
AND scout < 10
```

# Can have many queries in WITH

Another “view on the fly” syntax:

```
WITH Reds(bid, scout) AS  
(SELECT B.bid, COUNT (*)  
FROM Boats2 B, Reserves2 R  
WHERE R.bid = B.bid AND B.color = 'red'  
GROUP BY B.bid),
```

```
UnpopularReds AS  
SELECT bname, scout  
FROM Boats2 B, Reds  
WHERE Reds.bid=B.bid  
AND scout < 10
```

```
SELECT * FROM UnpopularReds;
```

# ARGMAX GROUP BY?

- The sailor with the highest rating per age

```
WITH maxratings(age, maxrating) AS  
(SELECT age, max(rating)  
FROM Sailors  
GROUP BY age)
```

```
SELECT S.*  
FROM Sailors S, maxratings m  
WHERE S.age = m.age  
AND S.rating = m.maxrating;
```

# Brief Detour: Null Values

- Field values are sometimes unknown
  - SQL provides a special value NULL for such situations.
  - Every data type can be NULL
- The presence of null complicates many issues. E.g.:
  - Selection predicates (WHERE)
  - Aggregation
- But NULLs comes naturally from Outer joins

# NULL in the WHERE clause

- Consider a tuple where rating IS NULL.

```
INSERT INTO sailors VALUES  
(11, 'Jack Sparrow', NULL, 35);
```

```
SELECT * FROM sailors  
WHERE rating > 8;
```

Is Jack Sparrow in the output?



# NULL in comparators

- Rule: (x op NULL) evaluates to ... NULL!

SELECT 100 = NULL;

SELECT 100 < NULL;

SELECT 100 >= NULL;

# Explicit NULL Checks

```
SELECT * FROM sailors WHERE rating IS NULL;
```

```
SELECT * FROM sailors WHERE rating IS NOT NULL;
```



# NULL at top of WHERE

- Rule: Do not output a tuple WHERE NULL

SELECT \* FROM sailors;

SELECT \* FROM sailors WHERE rating > 8;

SELECT \* FROM sailors WHERE rating <= 8;

# NULL in Boolean Logic

Three-valued logic:

NOT	T	F	N
	F	T	

AND	T	F	N
T	T	F	
F	F	F	
N			

OR	T	F	N
T	T	T	
F	T	F	
N			

SELECT \* FROM sailors WHERE rating > 8 AND TRUE;

SELECT \* FROM sailors WHERE rating > 8 OR TRUE;

SELECT \* FROM sailors WHERE NOT (rating > 8);

**General rule: NULL \*\*column values\*\* are ignored by aggregate functions**

# NULL in Boolean Logic

Three-valued logic:

NOT	T	F	N
	F	T	N

AND	T	F	N
T	T	F	N
F	F	F	F
N	N	F	N

OR	T	F	N
T	T	T	T
F	T	F	N
N	T	N	N

SELECT \* FROM sailors WHERE rating > 8 AND TRUE;

SELECT \* FROM sailors WHERE rating > 8 OR TRUE;

SELECT \* FROM sailors WHERE NOT (rating > 8);

**General rule: NULL \*\*column values\*\* are ignored by aggregate functions**

# NULL and Aggregation

```
SELECT count(*) FROM sailors;
```


```
SELECT count(rating) FROM sailors;
```

```
SELECT sum(rating) FROM sailors;
```

```
SELECT avg(rating) FROM sailors;
```

**General rule: NULL \*\*column values\*\* are ignored by aggregate functions**

# NULLs: Summary

- NULL op NULL is NULL
  - WHERE NULL: do not send to output
  - Boolean connectives: 3-valued logic
  - Aggregates ignore NULL-valued inputs
- 

# Testing SQL Queries

- SQL Fiddle pages we provide in this class will typically help you answer the questions in the worksheets and vitamins.
- But in real life:
  - not every database instance will reveal every bug in your query.
    - Eg: database instance without any rows in it!
  - Need to debug your queries
  - reasoning about them carefully
  - constructing test data.



# Tips for Generating Test Data

- Generate **random data**
  - e.g. using a service like [mockaroo.com](https://mockaroo.com)
- Try to construct data that could check for the following potential errors:
  - Incorrect output schema
  - Output may be missing rows from the correct answer (false negatives)
  - Output may contain incorrect rows (false positives)
  - Output may have the wrong number of duplicates.
  - Output may not be ordered properly.

# Summary

- You've now seen SQL—you are armed.
- A declarative language
  - Somebody has to translate to algorithms though...
  - The RDBMS implementor ... i.e. you!

# Summary Cont

- The data structures and algorithms that make SQL possible also power:
  - NoSQL, data mining, scalable ML, network routing...
  - A toolbox for scalable computing!
  - That fun begins next week
- We skirted questions of good database (schema) design
  - a topic we'll consider in greater depth later